Fully Paired Block Sparse Attention · blk64 → blk128

NVIDIA B20Z / SM100 | B=1 | Hq=Hkv=40 | Nq=Nkv=131072 | D=128 | BF16

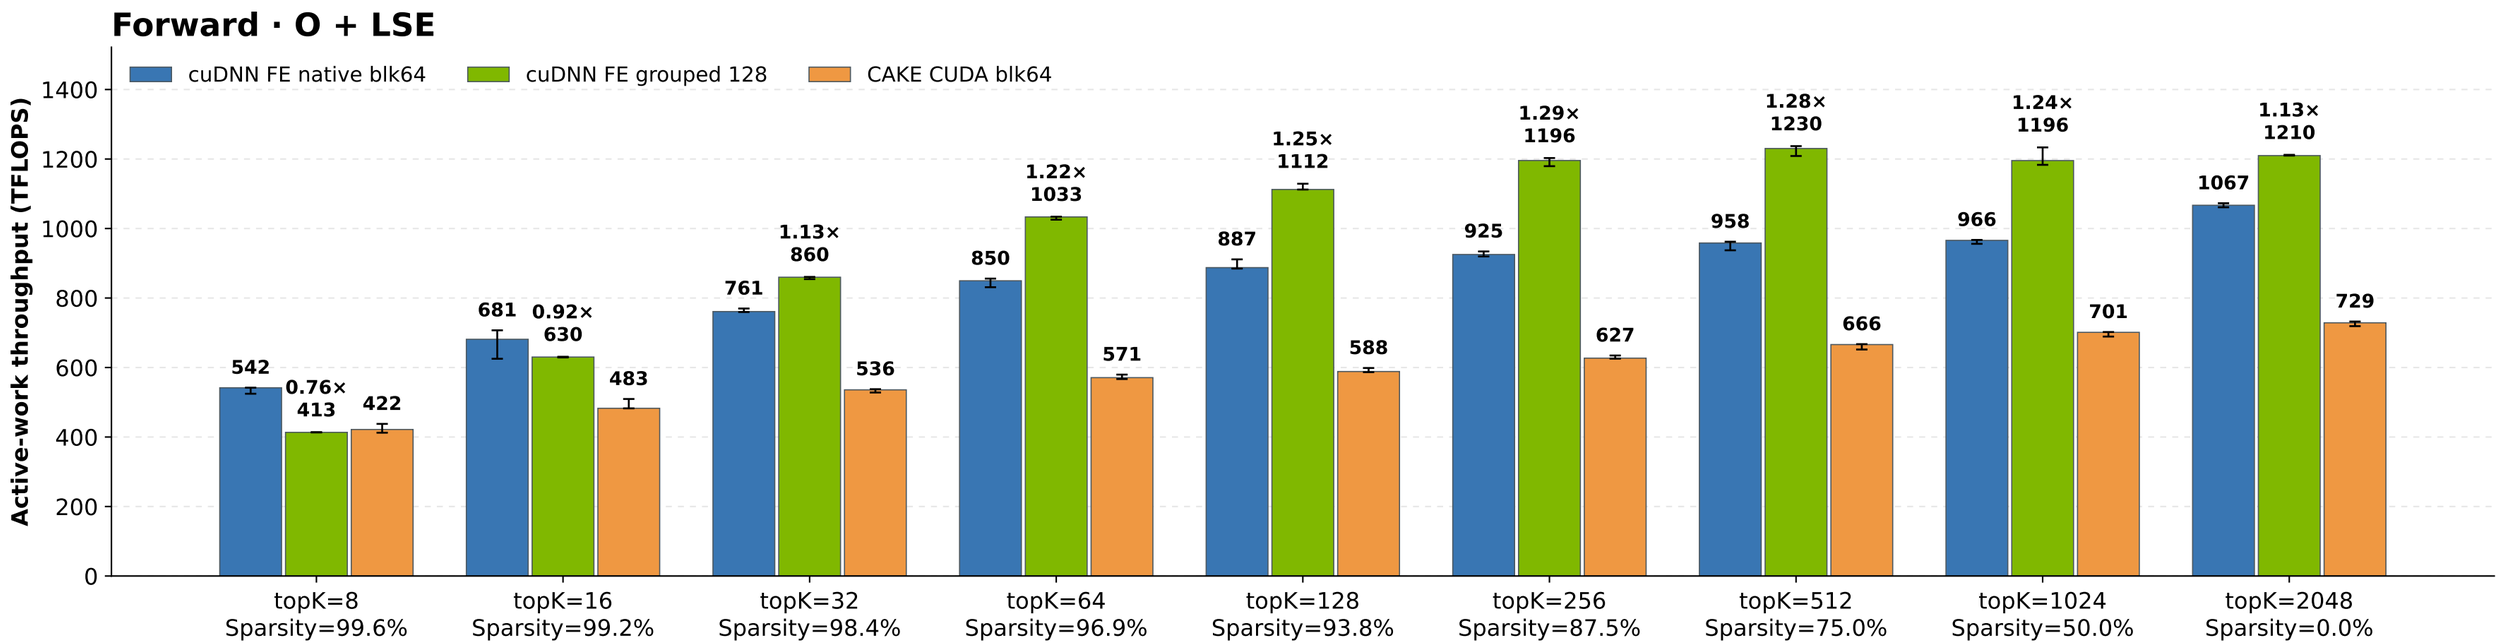

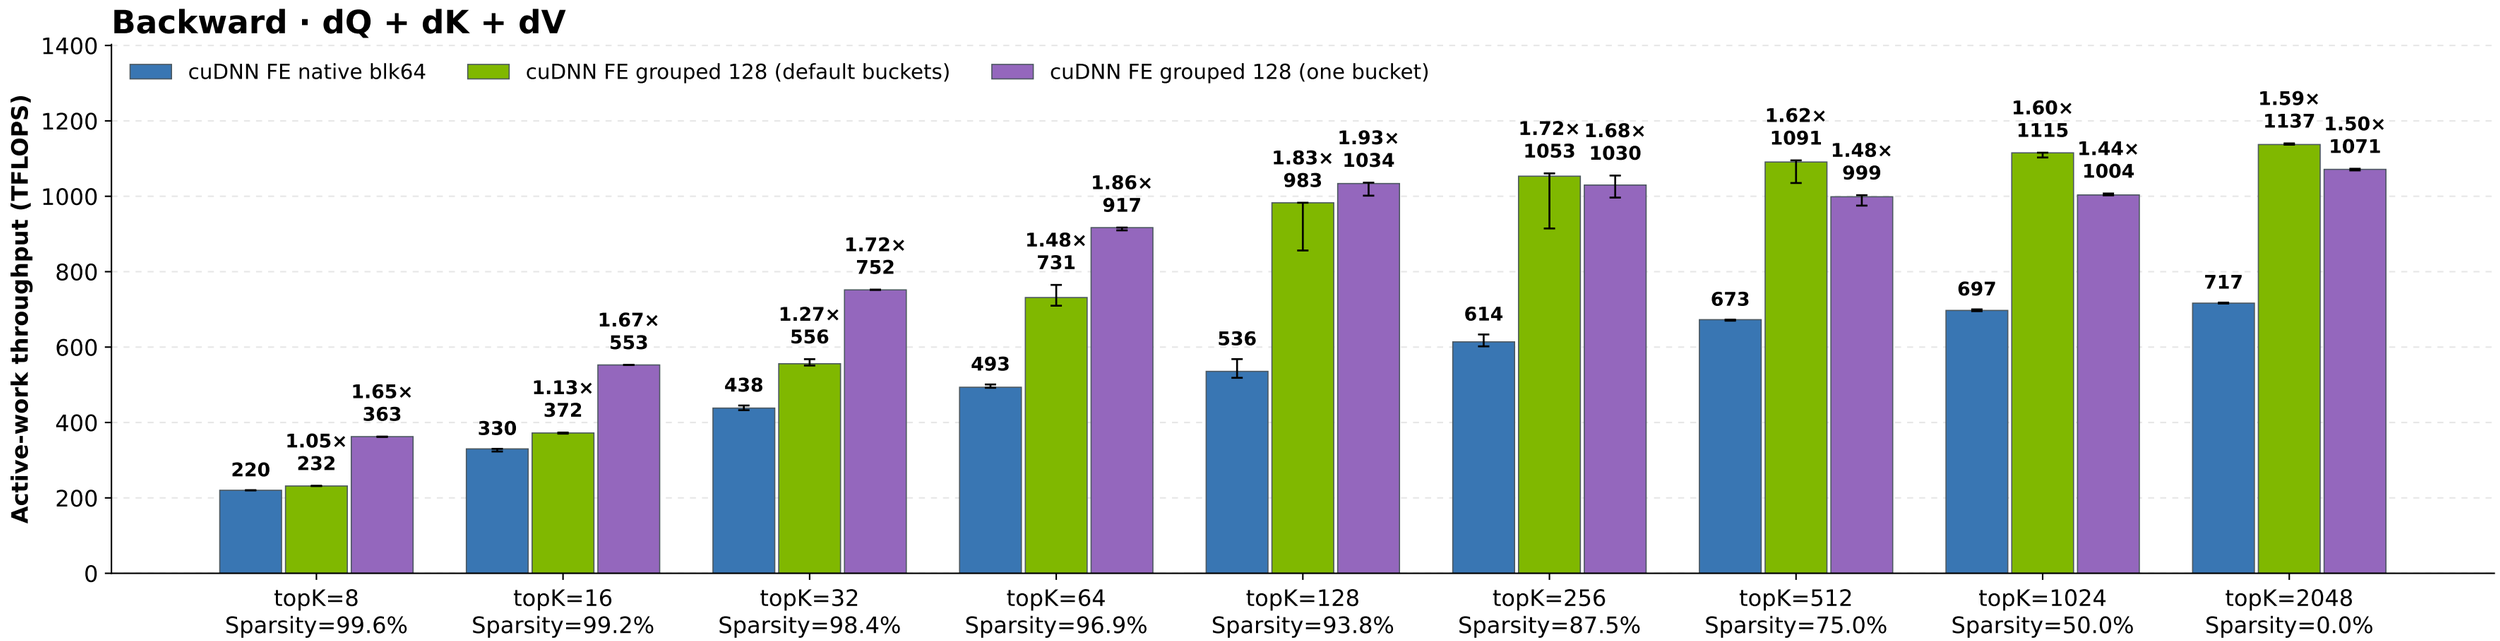

3 seeds × 7 randomized CUDA Graph rounds. Bars: median; whiskers: seed min–max (not confidence intervals). × = speedup vs native blk64.

Includes all grouped GPU permutations and backward reset/pre/post kernels. Excludes CPU preparation, routing, compilation and allocation. Synthetic 100% paired graph.